

A high-angle, black and white photograph of a massive concrete dam. The dam's surface is composed of large, rectangular blocks with visible vertical joints. A single person stands on the top edge of the dam, providing a sense of scale to the enormous structure. The sky is a uniform, dark grey.

INTRODUCTION TO

---

**EASYBUILD @ BI**

---

# PROBLEMS WITH SOFTWARE

- ▶ Installing and maintaining software from source requires a lot of time and effort
- ▶ Binary package systems focus on a single version compiled against a single set of compilers and libraries
- ▶ Installing new versions of software benefits very little from previous efforts
- ▶ We needed a way to automate installation of software for both system administrators and users

---

# BRIEF HISTORY

- ▶ Started in 2009, EasyBuild reached version 1.0, stable API and GPL license in 2012 (announced at SC'12)
- ▶ Developed originally by IT department of Ghent University in Belgium, member of Flemish Supercomputer Center, with monthly releases
- ▶ Most recent version 3.3.1 provides:
  - ▶ 1,262 different software recipes (easyconfigs) with 25 compiler toolchains
  - ▶ 187 software specific and 29 generic methods (easyblocks)

---

# WHAT IS EASYBUILD?

- ▶ Software build and installation framework, written in Python
- ▶ Supports different compilers and MPI libraries
- ▶ Very active and helpful community
- ▶ Features necessary to make it work for our environment at BI/VT were implemented by the core development team
- ▶ Software recipes are written by many people around the world

**SERIOUSLY,  
WHAT IS IT?**

---

# TERMINOLOGY (1 / 2)

- ▶ EasyBuild framework
  - ▶ core: Python modules and packages
  - ▶ provides supporting functionality for building and installing software
- ▶ easyblock
  - ▶ a Python Module that serves as a build script, 'plugin' for the EasyBuild framework
  - ▶ implements a (generic) software build/install procedure

---

# TERMINOLOGY (2 / 2)

- ▶ easyconfig file
  - ▶ (\*.eb): build specification, software name/version, compiler toolchain, source URL, dependency list
- ▶ (compiler) toolchain
  - ▶ set of compilers with accompanying libraries
- ▶ extensions
  - ▶ additional libraries/packages/modules for a particular application (e.g., Python, R )

---

# HOW DOES IT WORK?

- ▶ When EasyBuild is executed to follow an `easyconfig` recipe:
- ▶ `easyconfig` is evaluated for dependencies, including toolchains. Those are installed first, if missing.
- ▶ Source files are downloaded and unpacked
- ▶ Relevant configure, build, test and install steps are taken
- ▶ Each installation is verified, and if successful a new set of environment modules are created with full dependency information



## Sample easyconfig: ABySS/ABySS-1.9.0-foss-2016a.eb

```
# easy block selection
easyblock = 'ConfigureMake'

# software name and version
name = 'ABySS'
version = '1.9.0'

# software metadata
homepage = 'http://www.bcgsc.ca/platform/bioinfo/software/'
description = """Assembly By Short Sequences - a de novo"""

# toolchain name and version
toolchain = {'name': 'foss', 'version': '2016a'}
toolchainopts = {'usempi': True}

# sources and patches
sources = [SOURCELOWER TAR GZ]
source_urls = ['https://github.com/abyss/releases/%(version)s/']

# dependencies
dependencies = [
    ('Boost', '1.60.0'),
    ('sparsehash', '2.0.2'),
    ('SQLite', '3.9.2'),
]

# sanity checks
sanity_check_paths = {
    'files': ["bin/ABYSS", "bin/ABYSS-P"],
    'dirs': []
}

moduleclass = 'bio'
```

---

# WHAT'S DIFFERENT AT BI?

- ▶ Environment modules on each cluster pre-set variables for:
  - ▶ HPC administrators, managing globally installed software
  - ▶ Users, managing their local installations
- ▶ Users' environment allows access to ALL global software, and build locally any additional/missing items
  - ▶ Users can create their own easyconfigs, or override global ones
- ▶ Directory structure takes cluster name and architecture into account
  - ▶ a single home directory or /apps dir can be shared between many clusters

**HOW DO WE USE  
EASYBUILD?**

## Accessing already available software

```
$ ssh discovery1.bi.vt.edu
```

```
$ module list
```

```
Currently Loaded Modulefiles:
```

- 1) gcc/6.1.0
- 2) slurm/16.05.8
- 3) site/discovery/easybuild/setup

```
$ module load Valgrind/3.11.0-foss-2016a
```

```
$ module list
```

```
Currently Loaded Modulefiles:
```

- 1) gcc/6.1.0
- 2) slurm/16.05.8
- 3) site/discovery/easybuild/setup
- 4) GCCcore/4.9.3
- 5) binutils/2.25-GCCcore-4.9.3
- 6) GCC/4.9.3-2.25
- 7) numactl/2.0.11-GCC-4.9.3-2.25
- 8) hwloc/1.11.2-GCC-4.9.3-2.25
- 9) OpenMPI/1.10.2-GCC-4.9.3-2.25
- 10) OpenBLAS/0.2.15-GCC-4.9.3-2.25-LAPACK-3.6.0
- 11) gomp/2016a
- 12) FFTW/3.3.4-gomp-2016a
- 13) ScaLAPACK/2.0.2-gomp-2016a-OpenBLAS-0.2.15-LAPACK-3.6.0
- 14) foss/2016a
- 15) Valgrind/3.11.0-foss-2016a

```
$ module load <tab><tab>
```

```
Display all 531 possibilities? (y or n)
```

## Building new software (1)

```
# do we have 'pbzip2' installed?
```

```
dom@discovery1 ~> module av 2>&1 | grep -i pbzip
```

```
dom@discovery1 ~>
```

```
# let's load EasyBuild
```

```
dom@discovery1 ~> module load EasyBuild
```

```
# time to search for khmer
```

```
dom@discovery1 ~> eb -S pbzip2
```

```
CFG1=/apps/easybuild/ebfiles_repo-vbi/generic/p/PBZIP2
```

```
CFG2=/apps/easybuild/ebfiles_repo-vbi/generic/p/PBZIP2
```

```
CFG3=/apps/easybuild/software_discovery-sandy_bridge/EasyBuild/3.3.1/
```

```
lib/python2.7/site-packages/easybuild_easyconfigs-3.3.1-py2.7.egg/
```

```
easybuild/easyconfigs/p/PBZIP2
```

```
* $CFG1/PBZIP2-1.1.12-foss-2016a.eb
```

```
* $CFG1/PBZIP2-1.1.8-foss-2016a.eb
```

```
* $CFG2/PBZIP2-1.1.12-foss-2016a.eb
```

```
* $CFG2/PBZIP2-1.1.8-foss-2016a.eb
```

```
* $CFG3/PBZIP2-1.1.8-goolf-1.4.10.eb
```

```
* $CFG3/PBZIP2-1.1.8-ictce-6.2.5.eb
```

```
# let's see what's required to build it
```

```
dom@discovery1 ~> eb PBZIP2-1.1.12-foss-2016a.eb -Dr
```

```
Dry run: printing build status of easyconfigs and dependencies
```

```
CFG=/apps/easybuild/ebfiles_repo
```

```
* [x] $CFG/discovery-sandy_bridge/M4/M4-1.4.17.eb (module: M4/1.4.17)
```

```
[...]
```

```
* [ ] $CFG/vbi/generic/p/PBZIP2/PBZIP2-1.1.12-foss-2016a.eb (module:  
PBZIP2/1.1.8-foss-2016a)
```

## Building new software (2)

```
dom@discovery1 ~> eb PBZIP2-1.1.12-foss-2016a.eb --robot
```

```
== temporary log file in case of crash /tmp/eb-PSMVAS/easybuild-OiCo6M.log
== resolving dependencies ...
== processing EasyBuild easyconfig /apps/easybuild/ebfiles_repo-vbi/
generic/p/PBZIP2/PBZIP2-1.1.12-foss-2016a.eb
== building and installing PBZIP2/1.1.12-foss-2016a...
== fetching files...
== creating build dir, resetting environment...
== unpacking...
== patching...
== preparing...
== configuring...
== building...
== testing...
== installing...
== taking care of extensions...
== postprocessing...
== sanity checking...
== cleaning up...
== creating module...
== permissions...
== packaging...
== COMPLETED: Installation ended successfully
== Results of the build can be found in the log file(s) /home/dom/
easybuild/software/discovery-sandy_bridge/PBZIP2/1.1.12-foss-2016a/
easybuild/easybui
57.log
== Build succeeded for 1 out of 1
== Temporary log file(s) /tmp/eb-PSMVAS/easybuild-OiCo6M.log* have been
removed.
== Temporary directory /tmp/eb-PSMVAS has been removed.
```

## Building new software (3)

```
dom@discovery1 ~> module av 2>&1 | grep -i pbzip2
PBZIP2/1.1.12-foss-2016a
```

```
dom@discovery1 ~> module load PBZIP2/1.1.12-foss-2016a
```

```
dom@discovery1 ~> module list
Currently Loaded Modulefiles:
  1) gcc/6.1.0
  [...]
 17) PBZIP2/1.1.8-foss-2016a
```

```
dom@discovery1 ~> which pbzip2
~/easybuild/software/discovery-sandy_bridge/PBZIP2/1.1.12-foss-2016a/bin/pbzip2
```

```
dom@discovery1 ~> module av 2>&1 | less
```

```
[...]
----- /home/dom/easybuild/modules/discovery-
sandy bridge/all -----
PBZIP2/1.1.12-foss-2016a
```

```
----- /apps/easybuild/modules/discovery-
sandy bridge/all -----
ABYSS/1.9.0-foss-2016a
[...]
```

# WHAT COMES OUT?

- ▶ `~/easybuild/software/<cluster>-<arch>/<app_name>/<app_version>/`
  - ▶ full application
  - ▶ full build and install log
  - ▶ markdown report of the configure options, build, environment and install
- ▶ `~/easybuild/sources/<app_name_first_letter>/<app_name>`
  - ▶ source/installer
- ▶ `~/easybuild/modules/<cluster>-<arch>/all/<app_name>/<app_version>`
  - ▶ environment module for the application
- ▶ `~/easybuild/ebfiles_repo/<cluster>-<arch>/<app_name>/<app_name>-<app_version>-<toolchain_name-version>.eb`
  - ▶ new easyconfig, includes build stats



**IS THAT ALL?**

## Building new software - reusing existing easyconfigs

```
# installing different version from the same easyconfig
```

```
dom@discovery1 ~> eb PBZIP2-1.1.12-foss-2016a.eb --try-software-version=1.1.13
```

```
== temporary log file in case of crash /tmp/eb-wusUzW/easybuild-Eo5LXC.log
== resolving dependencies ...
== processing EasyBuild easyconfig /tmp/eb-wusUzW/tweaked_easyconfigs/
PBZIP2-1.1.13-foss-2016a.eb
== building and installing PBZIP2/1.1.13-foss-2016a...
== fetching files...
== creating build dir, resetting environment...
== unpacking...
== patching...
== preparing...
== configuring...
== building...
== testing...
== installing...
== taking care of extensions...
== postprocessing...
== sanity checking...
== cleaning up...
== creating module...
== permissions...
== packaging...
== COMPLETED: Installation ended successfully
== Results of the build can be found in the log file(s) /home/dom/easybuild/
software/discovery-sandy bridge/PBZIP2/1.1.13-foss-2016a/easybuild/easybuild-
PBZIP2-1.1.13-20170910.2I3532.log
== Build succeeded for 1 out of 1
== Temporary log file(s) /tmp/eb-wusUzW/easybuild-Eo5LXC.log* have been removed.
== Temporary directory /tmp/eb-wusUzW has been removed.
```

# IMPORTANT BITS

---

# HOW TO GET NEW SOFTWARE

- ▶ First, check if it's already installed
  - ▶ `module av 2>&1 | grep -i <app>`
- ▶ Then, see if easyconfig is available
  - ▶ `module load EasyBuild`
  - ▶ `eb -S <app>`
  - ▶ `eb <easyconfig.eb> -robot`
- ▶ If you need a newer/different version, try:
  - ▶ `eb <easyconfig.eb> -robot -try-software-version=<x.y>`

---

# IMPORTANT BITS

- ▶ Most easyconfigs target either FOSS or INTEL toolchains, we have standardized on FOSS:
  - ▶ foss/2016a
  - ▶ foss/2016b
  - ▶ foss/2017a
- ▶ foss toolchain comprises of:
  - ▶ compiler – GCC
  - ▶ MPI implementation – OpenMPI
  - ▶ libraries – OpenBLAS/LAPACK, ScaLAPACK(/BLACS), FFTW

---

# CREATING NEW EASYCONFIGS

- ▶ Find an existing similar easyconfig from foss/2016a or foss/2016b toolchain
- ▶ Share the easyconfig between the group, or send to [helpdesk@rt.bi.vt.edu](mailto:helpdesk@rt.bi.vt.edu) and I'll include it in the global installs
- ▶ Examples, examples, examples! I constantly browse existing easyconfigs to get ideas
- ▶ Be sure to use proper naming scheme
  - ▶ pigz-2.3.3-foss-2016b.eb
- ▶ Place them in either location:
  - ▶ ~/easybuild/ebfiles\_repo/generic/
  - ▶ ~/easybuild/ebfiles\_repo/<cluster>-<arch>/

# NOTES

- ▶ Always start with a clean module environment, do not preload modules in your `~/.bashrc` or `~/.bash_profile`
- ▶ In your slurm/PBS scripts load module only for the application you are trying to run. EasyBuild generated modules have all the dependency requirements needed
- ▶ When re-using easyconfigs from different toolchains, match dependency software versions to the new toolchain
  - ▶ `module av 2>&1 | grep foss-2016b`

---

# GETTING HELP

- ▶ EasyBuild documentation is excellent
  - ▶ <https://easybuild.readthedocs.io/>
  - ▶ `eb --help`
- ▶ EasyBuild developers and users are available via:
  - ▶ mailing list
  - ▶ IRC channel
  - ▶ Slack channel
- ▶ Contact us via the RT system:  
[helpdesk@rt.bi.vt.edu](mailto:helpdesk@rt.bi.vt.edu)



---

# BI EASYCONFIGS AND ENVIRONMENT SCRIPTS

- ▶ Complete and up-to-date list of all easyconfigs and modules installed on BI clusters
- ▶ <https://devlab.vbi.vt.edu/HPC/easybuild>
- ▶ [https://github.com/dominikborkowski/biocomplexity\\_easybuild](https://github.com/dominikborkowski/biocomplexity_easybuild)
- ▶ includes entry level documentation and design

**DEMO TIME!**